

CS 61C

MR, Spark, WSC, RAID, ECC

Summer 2020

Discussion 14: August 10, 2020

1 Pre-Check

This section is designed as a conceptual check for you to determine if you conceptually understand and have any misconceptions about this topic. Please answer true/false to the following questions, and include an explanation:

1.1 MapReduce is more general than Spark since it is lower level.

1.2 The higher the PUE the more efficient the datacenter is.

1.3 Hamming codes can detect any type of data corruption.

1.4 All RAID levels improve reliability.

2 Hamming ECC

Recall the basic structure of a Hamming code. We start out with some bitstring, and then add parity bits at the indices that are powers of two (1, 2, 8, etc.). We don't assign values to these parity bits yet. **Note that the indexing convention used for Hamming ECC is different from what you are familiar with.** In particular, the 1 index represents the MSB, and we index from left-to-right. The i th parity bit $P\{i\}$ covers the bits in the new bitstring where the *index* of the bit under the aforementioned convention, j , has a 1 at the same position as i when represented as binary. For instance, 4 is 0b100 in binary. The integers j that have a 1 in the same position when represented in binary are 4, 5, 6, 7, 12, 13, etc. Therefore, $P4$ covers the bits at indices 4, 5, 6, 7, 12, 13, etc. A visual representation of this is:

Bit position		1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	
Encoded data bits		p1	p2	d1	p4	d2	d3	d4	p8	d5	d6	d7	d8	d9	d10	d11	p16	d12	d13	d14	d15	
Parity bit coverage	p1	X		X		X		X		X		X		X		X		X		X		
	p2		X	X			X	X			X	X			X	X			X	X		...
	p4				X	X	X	X					X	X	X	X					X	
	p8								X	X	X	X	X	X	X	X						
	p16																X	X	X	X	X	

Source: https://en.wikipedia.org/wiki/Hamming_code

2.1 How many bits do we need to add to 0011₂ to allow single error correction?

2 *MR, Spark, WSC, RAID, ECC*

- 2.2 Which locations in 0011_2 would parity bits be included?
- 2.3 Which bits does each parity bit cover in 0011_2 ?
- 2.4 Write the completed coded representation for 0011_2 to enable single error correction. Assume that we set the parity bits so that the bits they cover have even parity.
- 2.5 How can we enable an additional double error detection on top of this?
- 2.6 Find the original bits given the following SEC Hamming Code: 0110111_2 . Again, assume that the parity bits are set so that the bits they cover have even parity.
- 2.7 Find the original bits given the following SEC Hamming Code: 1001000_2

3 RAID

- 3.1 Fill out the following table:

	Configuration	Pro/Good for	Con/Bad for
RAID 0			
RAID 1			
RAID 2			
RAID 3			
RAID 4			

RAID 5			
--------	--	--	--

4 MapReduce

For each problem below, write pseudocode to complete the implementations. Tips:

- The input to each MapReduce job is given by the signature of `map()`.
- `emit(key k, value v)` outputs the key-value pair `(k, v)`.
- `for var in list` can be used to iterate through `Iterables` or you can call the `hasNext()` and `next()` functions.
- Usable data types: `int`, `float`, `String`. You may also use lists and custom data types composed of the aforementioned types.
- `intersection(list1, list2)` returns a list of the common elements of `list1`, `list2`.

- 4.1 Given a set of coins and each coin's owner in the form of a list of `CoinPairs`, compute the number of coins of each denomination that a person has.

`CoinPair`:

```
String person
String coinType
```

```
1 map(CoinPair pair):                                1 reduce(_____, _____):
```

- 4.2 Using the output of the first MapReduce, compute each person's amount of money. `valueOfCoin(String coinType)` returns a float corresponding to the dollar value of the coin.

```
1 map(tuple<CoinPair, int> output):                    1 reduce(_____, _____):
```

4 *MR, Spark, WSC, RAID, ECC*

5 Spark

Resilient Distributed Datasets (RDD) are the primary abstraction of a distributed collection of items

Transforms $RDD \rightarrow RDD$

map(f) Return a new transformed item formed by calling f on a source element.

flatMap(f) Similar to map, but each input item can be mapped to 0 or more output items (so f should return a sequence rather than a single item).

reduceByKey(f) When called on a dataset of (K, V) pairs, returns a dataset of (K, V) pairs where the values for each key are aggregated using the given reduce function f , which must be of type $(V, V) \rightarrow V$.

Actions $RDD \rightarrow Value$

reduce(f) Aggregate the elements of the dataset *regardless of keys* using a function f .

Call `sc.parallelize(data)` to parallelize a Python collection, `data`.

- 5.1 Given a set of coins and each coin's owner, compute the number of coins of each denomination that a person has. Then, using the output of the first result, compute each person's amount of money. Assume `valueOfCoin(coinType)` is defined and returns the dollar value of the coin.

The type of `coinPairs` is a tuple of `(person, coinType)` pairs.

```
1 coinData = sc.parallelize(coinPairs)
```

- 5.2 Given a student's name and course taken, output their name and total GPA.

CourseData:

```
int courseID
float studentGrade // a number from 0-4
```

The type of `students` is a list of `(studentName, courseData)` pairs.

```
1 studentsData = sc.parallelize(students)
```

6 Warehouse-Scale Computing

Sources speculate Google has over 1 million servers. Assume each of the 1 million servers draw an average of 200W, the PUE is 1.5, and that Google pays an average of 6 cents per kilowatt-hour for datacenter electricity.

- 6.1

 Estimate Google's annual power bill for its datacenters.

- 6.2

 Google reduced the PUE of a 50,000-machine datacenter from 1.5 to 1.25 without decreasing the power supplied to the servers. What's the cost savings per year?